

Convolutional neuronal networks

Johannes Müller

With Material from

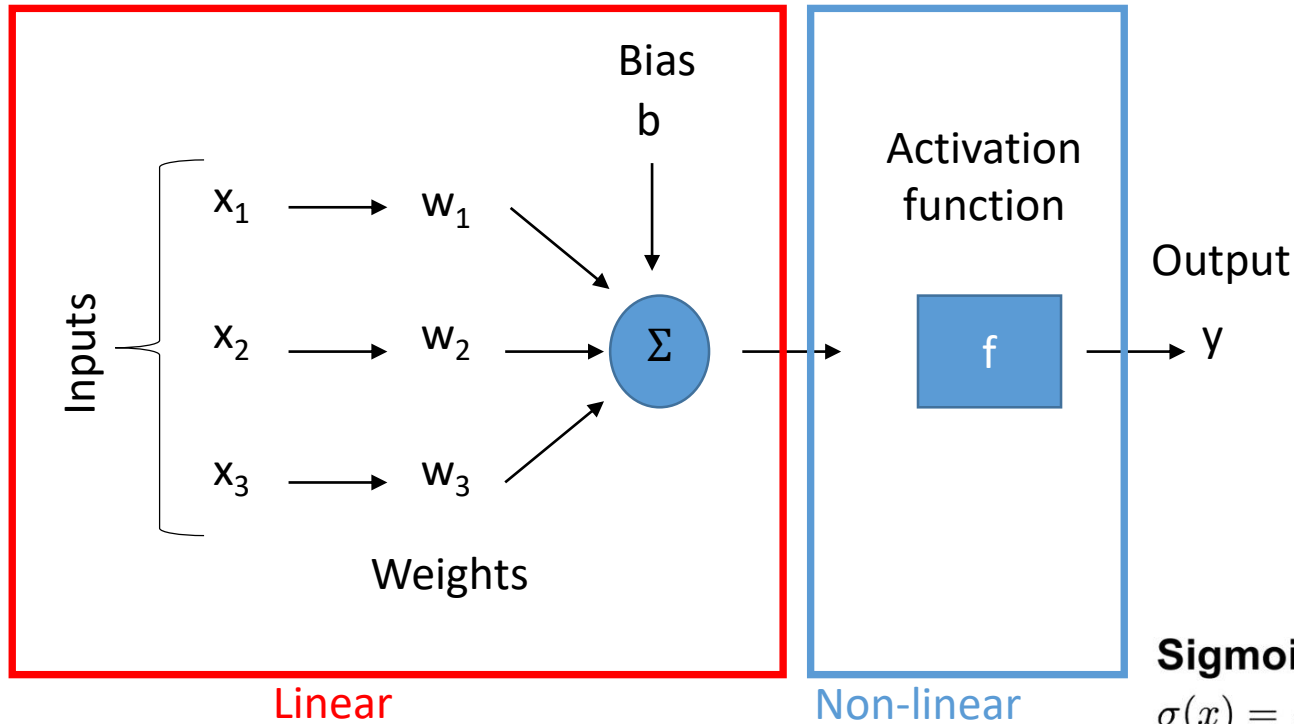
Robert Haase, PoL

Alex Krull, MPI CBG

Martin Weigert, EPFL Lausanne

Uwe Schmidt, MPI CBG

Ignacio Arganda-Carreras, Universidad del Pais Vasco



Single neuron output calculation

$$y = w_1x_1 + w_2x_2 + w_3x_3 + b = w^T x + b$$

For image data, the values x_1, x_2, \dots would be


Pixel intensities

Pixel coordinates

Filter kernel entries


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

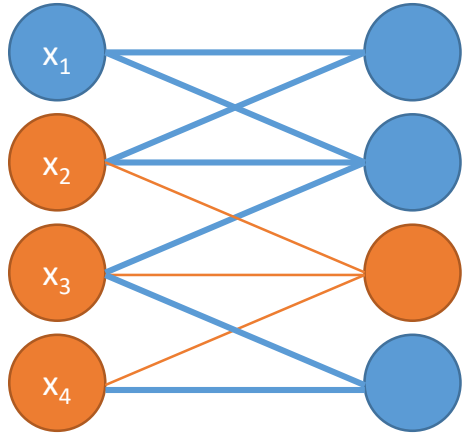
$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

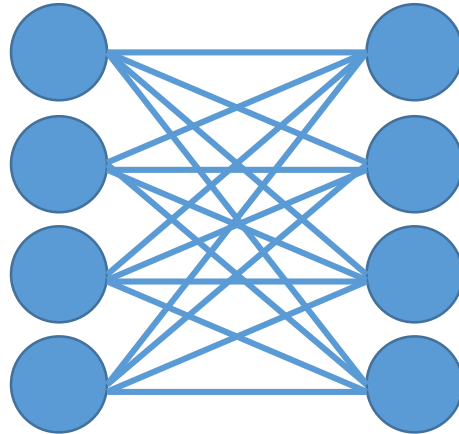
$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



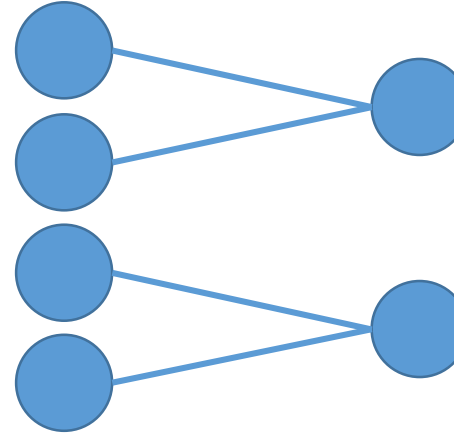
- Layers



Convolutional layer



Fully connected layer

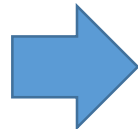


Pooling layer
("Max pool", "Average pool") Pooling maximal values

Previously:

Defined filter kernels

1/16	1/8	1/16
1/8	1/4	1/8
1/16	1/8	1/16




Now:

Undefined filter kernels

w_{11}	w_{12}	w_{13}
w_{21}	w_{22}	w_{23}
w_{31}	w_{32}	w_{33}

3	15	1	13
9	7	0	10
11	5	5	3
1	8	9	6



15	13
11	9


Averaging values

8.5	6.0
6.3	5.8

- Learning is an optimization problem
- Step 0: Initialize the network randomly
 - Weights
 - Bias
- Step 1: Forward pass the input through the network, get an initial prediction (Images 0...M)
- Step 2: Compare the output with the ground truth, compute the error (loss function)
 - The **loss function** can be freely defined.
 - Mean squared error:

$$\mathcal{L}(y, \hat{y}) = \frac{1}{M} \sum_{i=1}^M (\hat{y}_i - y_i)^2$$

- Step 3: Update weights



The loss function can be expanded from

$$\mathcal{L}(y, \hat{y}) = \frac{1}{M} \sum_{i=1}^M (\hat{y}_i - y_i)^2$$

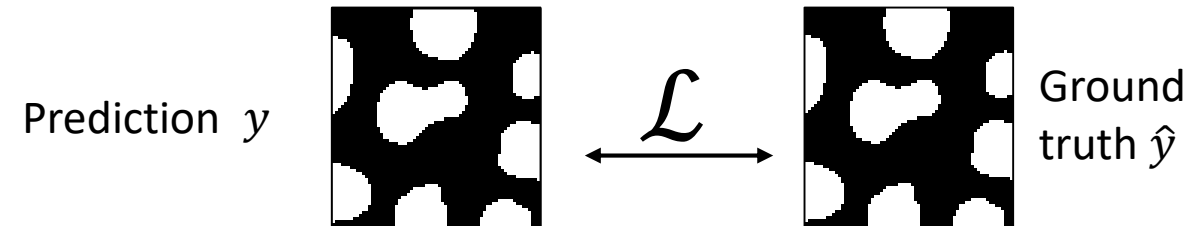
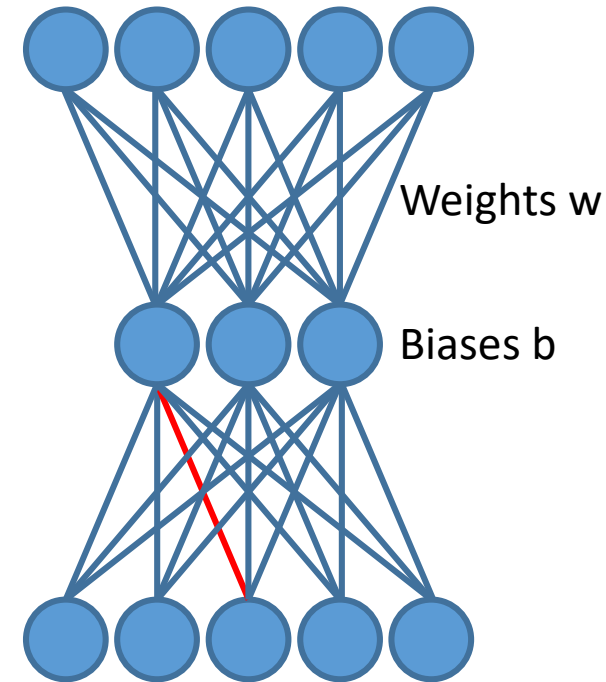
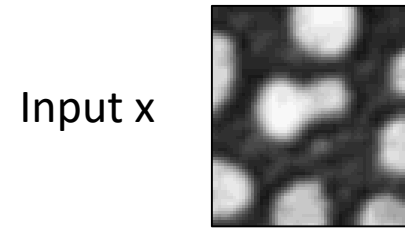
as the prediction depends on inputs x weights w and bias b

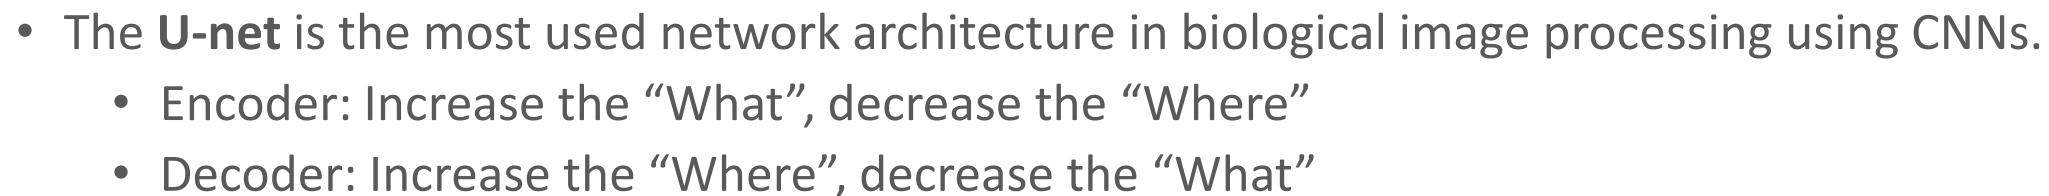

$$\mathcal{L}(\hat{y}, x, w) = \frac{1}{M} \sum_{i=1}^M (\hat{y}_i - (w^T x_i + b))^2$$

We can calculate derivatives with respect to w and b to find their optimal values


→ Derivatives tell us how to change w & b in order to improve the prediction

Repeat this n times, each time update weights w





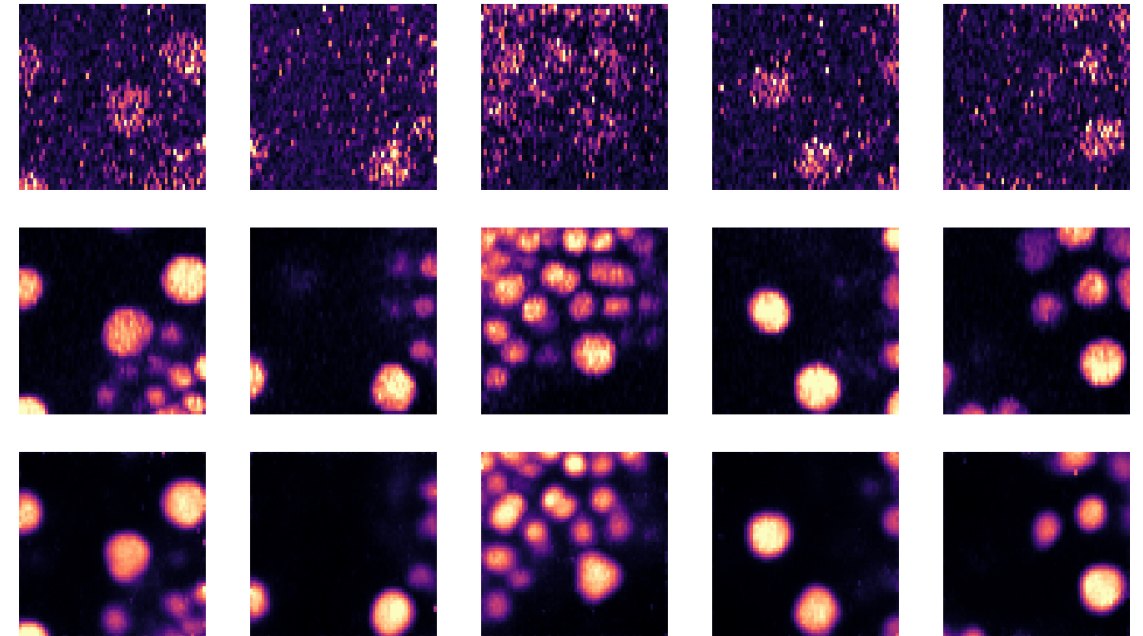
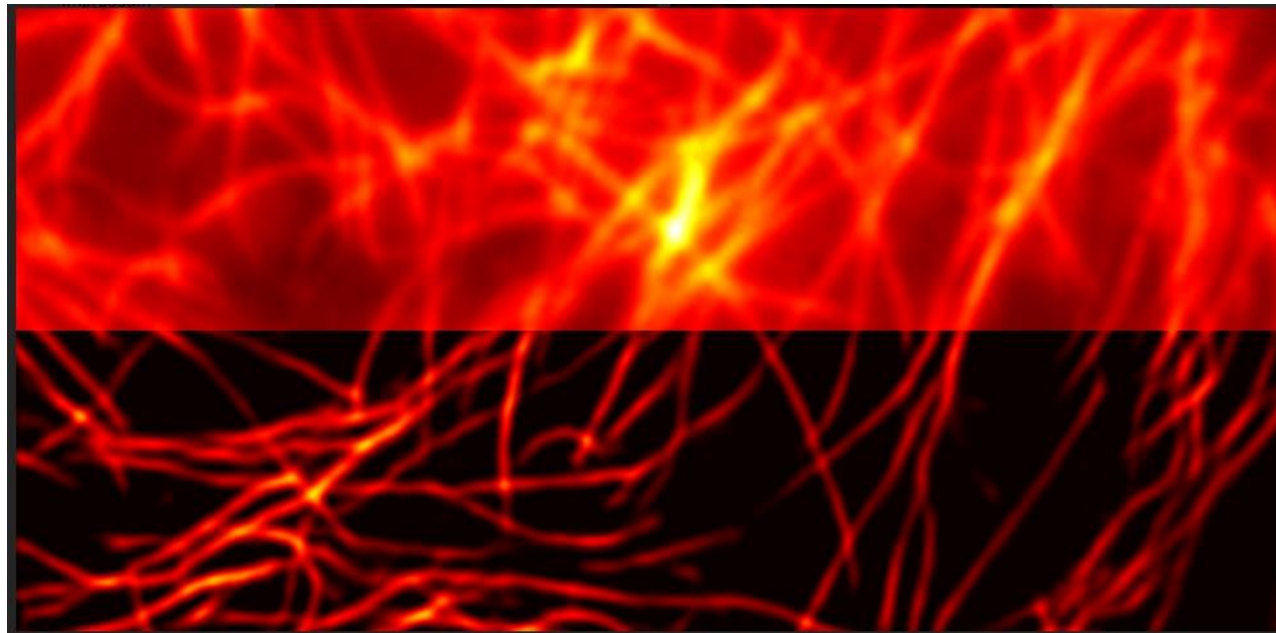
YOLO: You only look once



<https://github.com/ultralytics/yolov5>


- **CARE:** content-aware restoration
- Image acquisition of pairs of images: A high-quality and a low-quality image.
- Problem: Shot noise, Biology moves!
- Trained model only applicable to image data of the same conditions (biological system, microscope, etc)

5 example validation patches
top row: input (source), middle row: target (ground truth), bottom row: predicted from source





Raw data = signal + noise



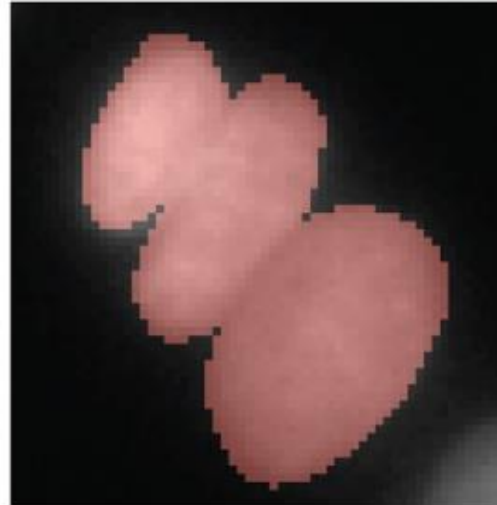
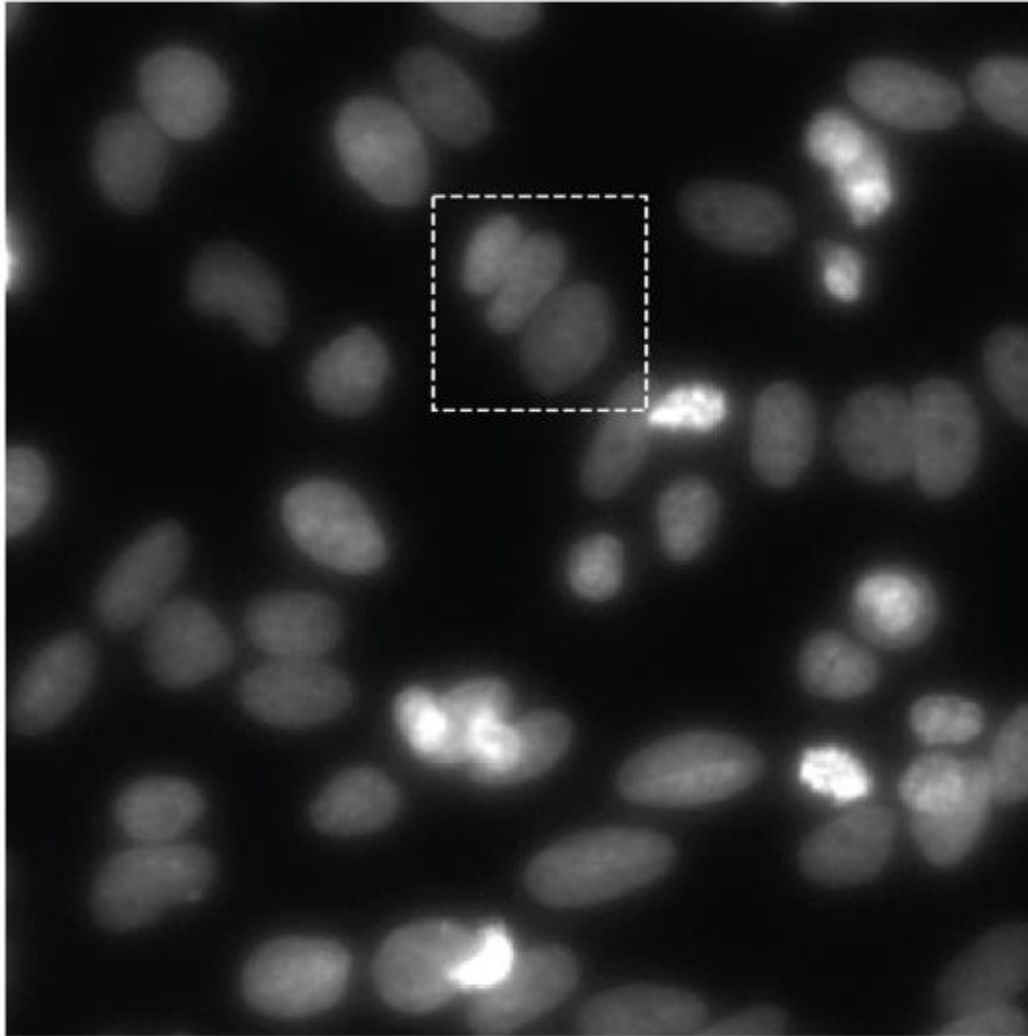
Strategy:

- Try to predict intensity of pixel y from surrounding pixels x
- CNN fails to predict noise component → N2V can only reproduce signal from the surroundings of y
- Only **random** noise can be removed, otherwise artifacts occur

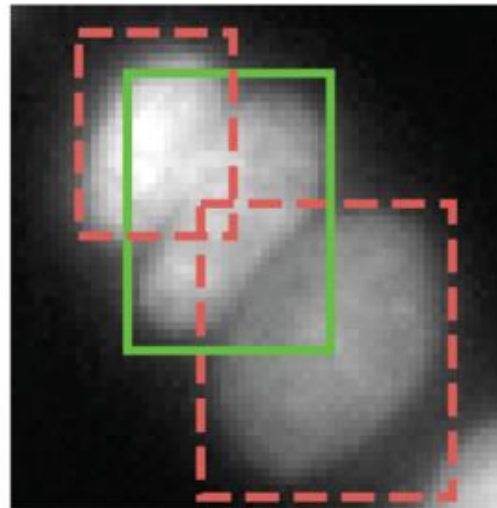
<https://github.com/juglab/n2v>

<https://forum.image.sc/t/n2v-artefacts-in-training-data/70686>

Noisy images + Crowded cells = Common source of segmentation errors



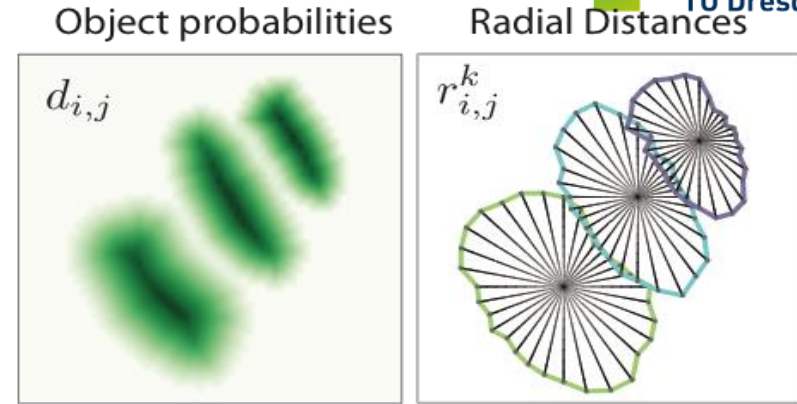
Dense Segmentation
(e.g. U-Net)




Bounding box based methods
(e.g. Mask-RCNN)

Strategy:

- Add additional information to prediction
- Member pixels of objects (nuclei) can be reached via a straight line from the center



NMS



Dense Polygon Prediction
(e.g. U-Net, ResNet)

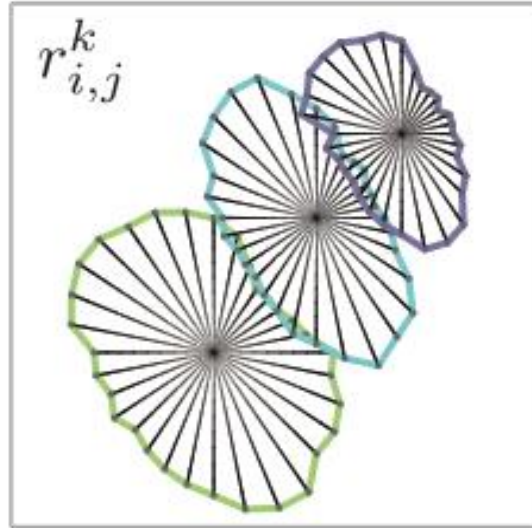
Polygon Selection
(Non-Maximum Suppression NMS)

Schmid et al., MICCAI (2018)

Object probabilities



Radial Distances



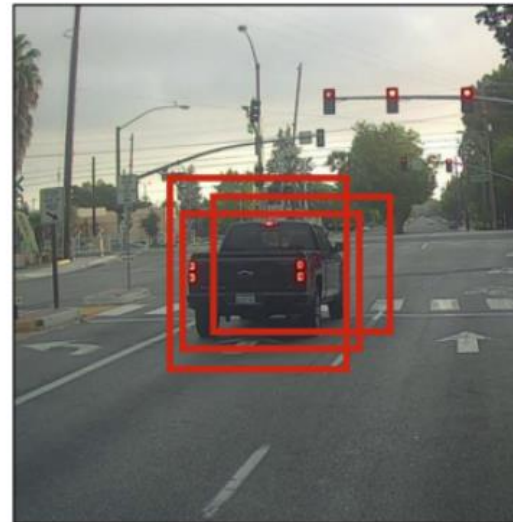
Problem:

- Multiple candidate points for nucleus center
- Overlapping instance predictions

Non-maximum-suppression (NMS):

- Intersection over Union (IoU) threshold τ determines „conservativeness“:
 - High τ : Objects tend to be considered as separate objects
 - Low τ : Objects tend to be considered as the same objects

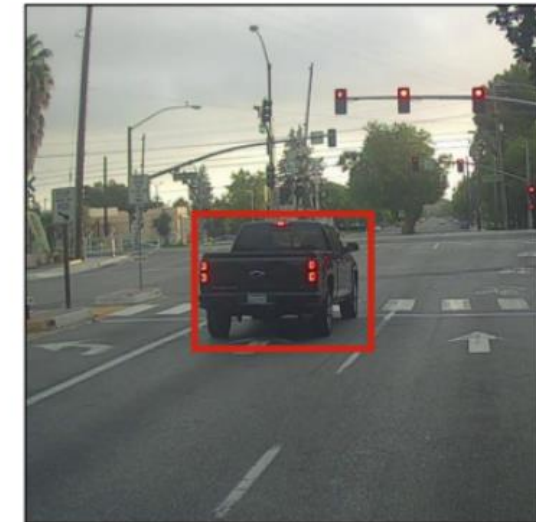
Before non-max suppression



Non-Max
Suppression



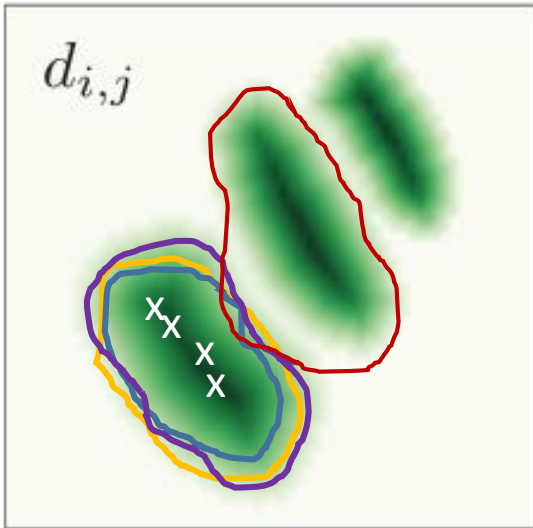
After non-max suppression



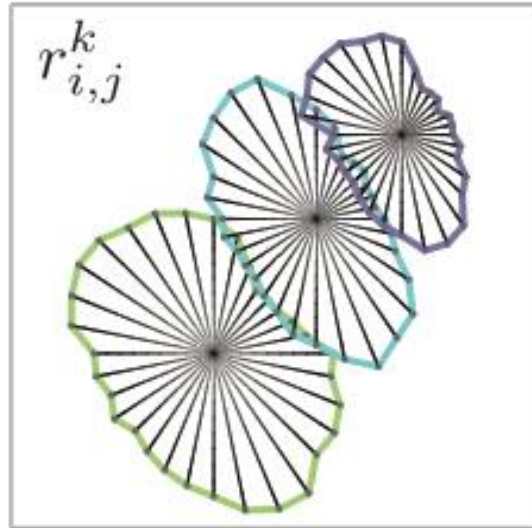
<https://towardsdatascience.com/non-maximum-suppression-nms-93ce178e177c>

Schmid et al., MICCAI (2018)

Object probabilities











Radial Distances

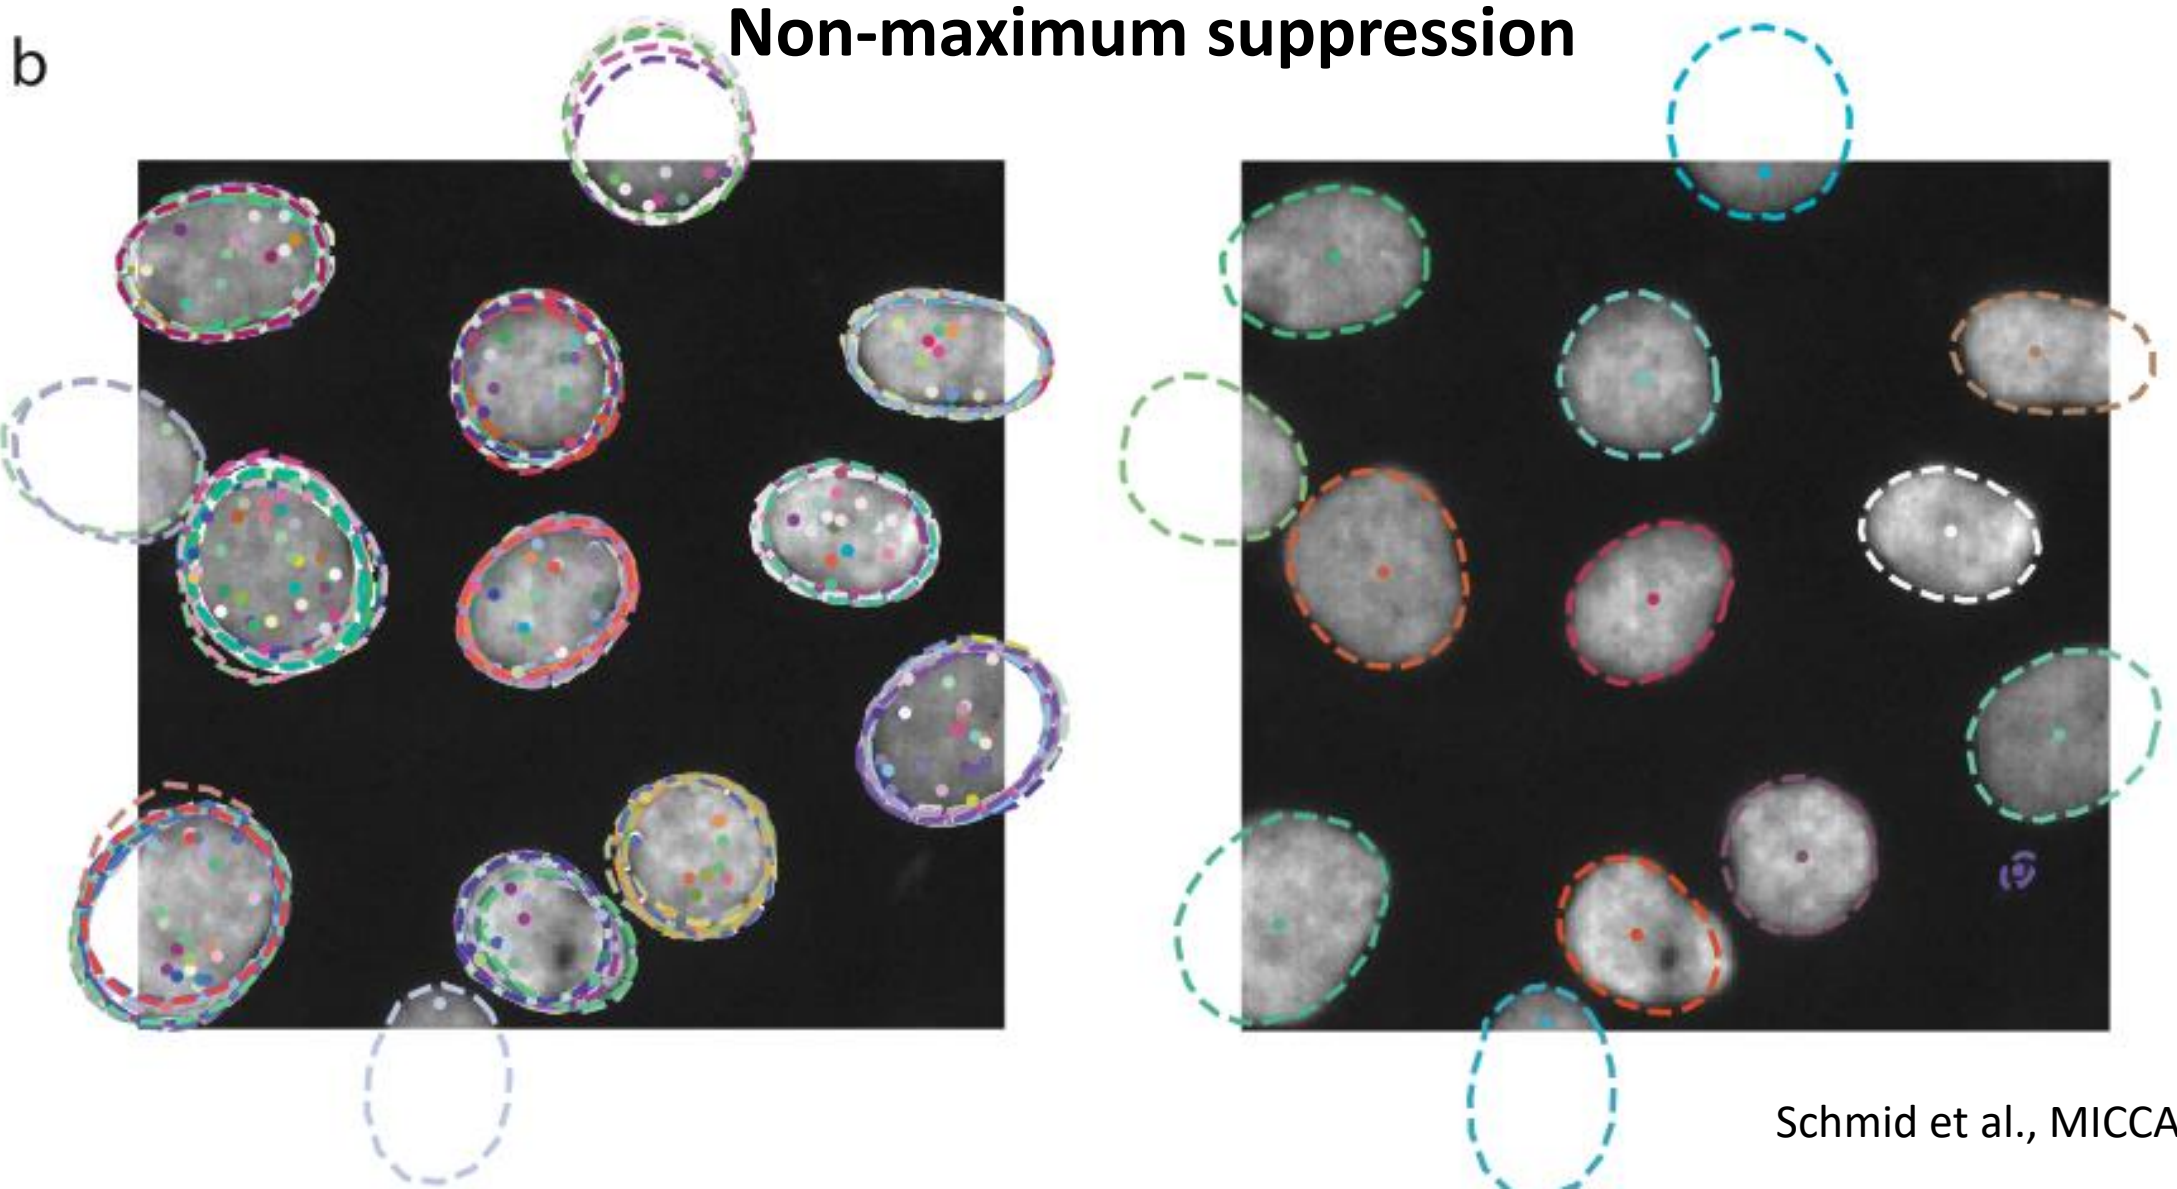


Non-maximum-suppression (NMS):

- Object probabilities: Probability that pixel belongs to class “nucleus”
- Multiple maxima lead to multiple possible polygons for the same nucleus

Algorithm:

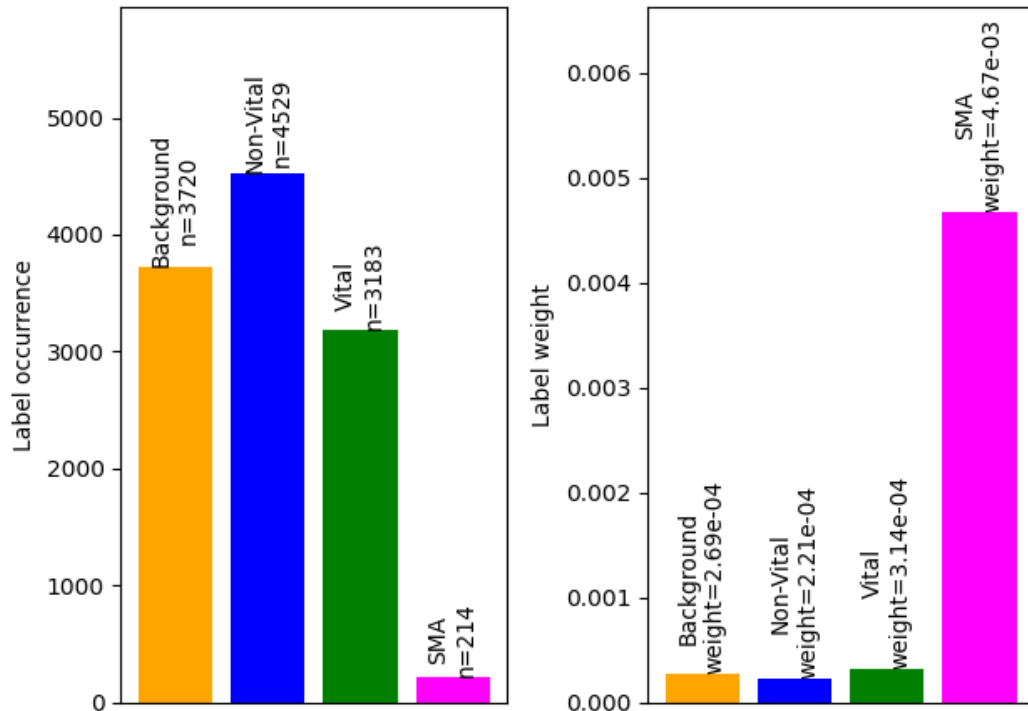
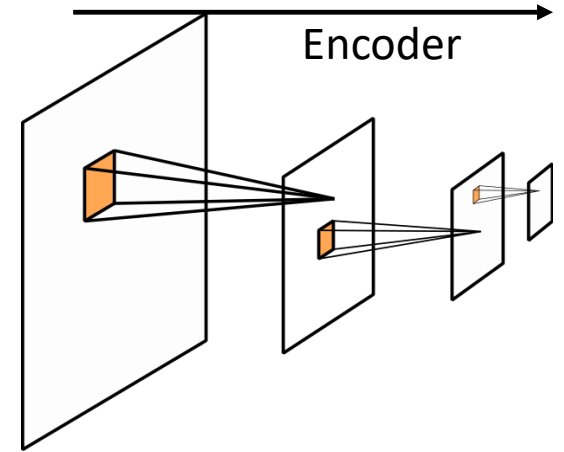
- Select polygon with highest object probability inside: 
- Look at other polygons: Is the overlap of  with  larger than threshold τ ?
 - Yes:  and  are actually the same object, drop 
 - No:  and  are separate nuclei
- Setting τ very high leads to many false positives!



Many prediction frameworks use UNets – similar weak points

Receptive field:

- Neurons in deeper layers can only “see” parts of the raw image
- Objects must be smaller than receptive field to be detectable



Unbalanced training data:

- Heterogeneous occurrence of labels in training data
- Rare events will not be caught because they don't harm accuracy much
- Weighted data sampling

Is the iPhone racist? Chinese users claim iPhoneX face recognition can't tell them apart

APPLE has come under fire following numerous complaints from Chinese users who claim the iPhone X face recognition can't tell them apart.

<https://www.news.com.au/>

Popular frameworks: <https://www.tensorflow.org/> , <https://www.pytorch.org/>

PyTorch

TensorFlow

Hardware requirements: Nvidia (CUDA-capable) graphics card (GPU)

Memory: GPU memory limits < RAM memory → Images are tiled and passed through the network in *batches*

Batch-processing:

Batch dimension

Channel dimension

→ For 2D multichannel image: $\text{data.shape} = [B, C, Y, X]$

→ Batch normalization: Data in batch is normalized from $[-\sigma, \sigma] \rightarrow [-1, 1]$

Higher batch size, smaller tiles → Better generalization of image statistics

Lower batch size, bigger tiles → Potentially larger receptive field


Train/test/validation split:

→ During training: Measure performance (“*loss*”) of network, **update weights**

→ During testing: Measure performance of updated network, **keep weights**

→ Validation: Measure performance in **unseen** validation dataset

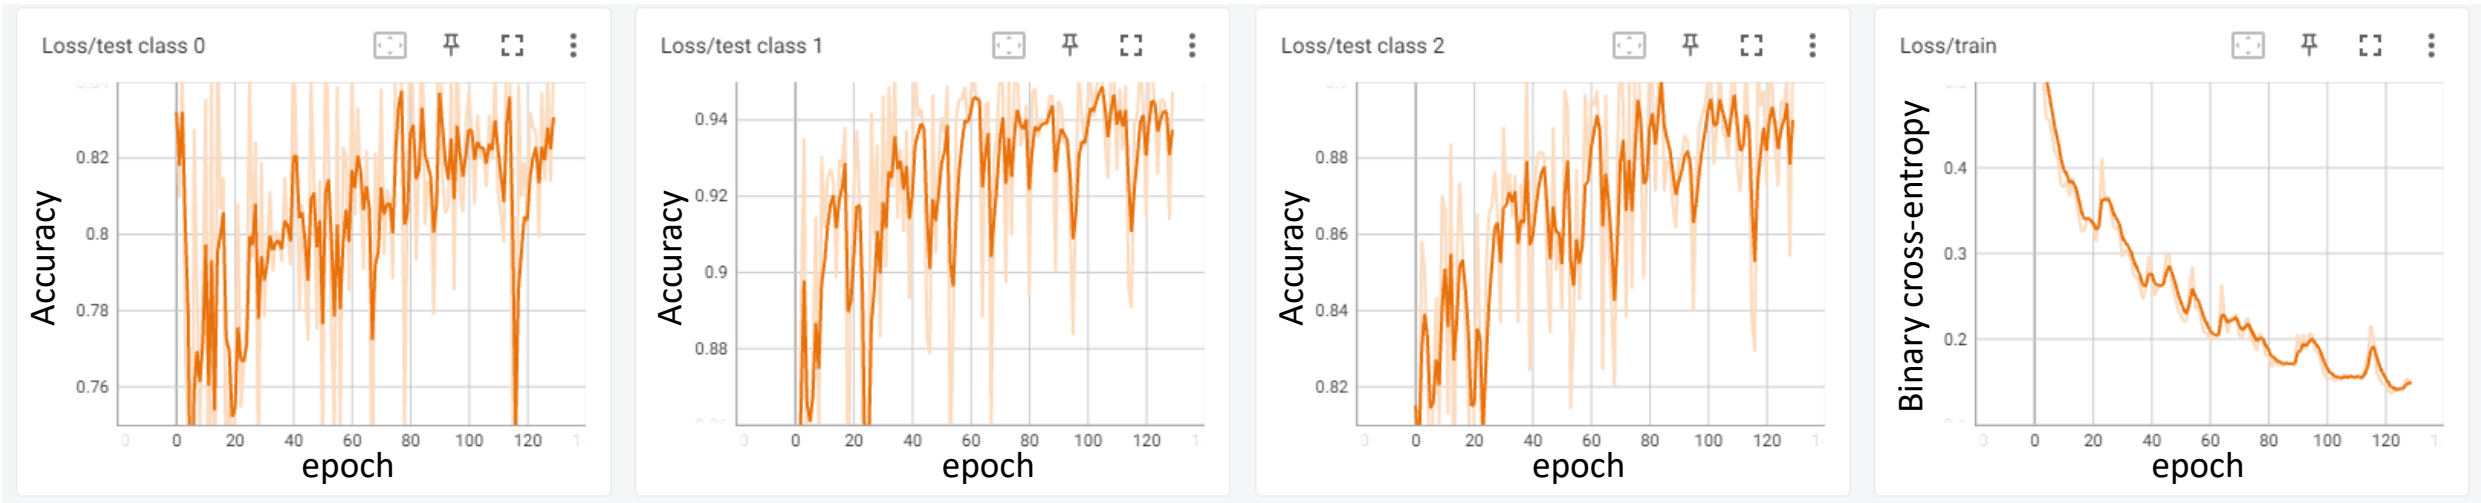
1 epoch



Tiles in batch

Tiles in batch

Training loss: Binary cross-entropy (Measure for information in image A contained in image B)

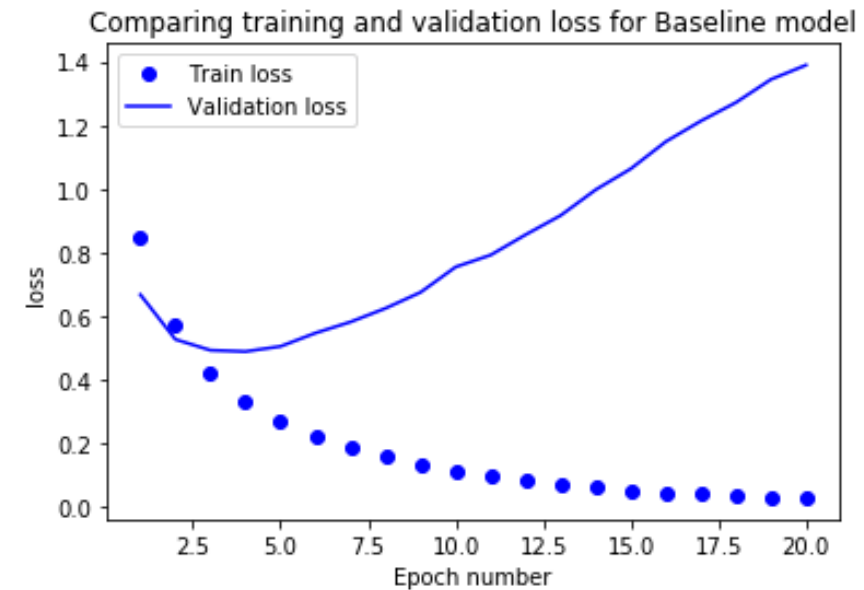



Test loss: Label-wise accuracy

→ We use different metrics in training/testing!

Overfitting:

- Network is learning things „by heart“
- Hint at this happening: Updated weights from training fail to perform well in test





0th-batch

```
data['image'][0].shape
```

```
torch.Size([3, 256, 256])
```


```
data['prediction'][0,0]
```

```
data['prediction'][0,1]
```



```
data['prediction'][0,2]
```

Stride:


- Network can not predict accurately at image edges
- We throw away a X-pixel wide margin at the image edge



Example



Prediction






What happened here?

Receptive field too small

I used a different resolution than during training

Overfitting



- With great power comes great responsibility: **Validate your models well!**
- Better data > better model